

Assignment #03

Yitong Wu

February 18, 2020

(a) Use your own synthesis model to synthesize $m = 1$ synthetic dataset for the CE sample.

```
library(readr)
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
library(ggplot2)
library(fastDummies)
knitr::opts_chunk$set(echo = TRUE)
def.chunk.hook <- knitr::knit_hooks$get("chunk")
knitr::knit_hooks$set(chunk = function(x, options) {
  x <- def.chunk.hook(x, options)
  ifelse(options$size != "normalsize", paste0("\n", options$size, "\n\n", x, "\n\n \\normalsize"), x)
})
require(runjags)

## Loading required package: runjags
require(coda)

## Loading required package: coda
data <- read.csv("CEdata.csv")
data$Rural = fastDummies::dummy_cols(data$UrbanRural[,names(fastDummies::dummy_cols(data$UrbanRural))])
data$Race_Black = fastDummies::dummy_cols(data$Race[,names(fastDummies::dummy_cols(data$Race)) == ".data_1"])
data$Race_NA = fastDummies::dummy_cols(data$Race[,names(fastDummies::dummy_cols(data$Race)) == ".data_2"])
data$Race_Asian = fastDummies::dummy_cols(data$Race[,names(fastDummies::dummy_cols(data$Race)) == ".data_3"])
data$Race_PI = fastDummies::dummy_cols(data$Race[,names(fastDummies::dummy_cols(data$Race)) == ".data_4"])
data$Race_M = fastDummies::dummy_cols(data$Race[,names(fastDummies::dummy_cols(data$Race)) == ".data_6"])
data$logInc <- log(data$Income)

modelString <- "
model {
  ## sampling
  for (i in 1:N){
    y[i] ~ dnorm(beta0 + beta1*x_rural[i] +
      beta2*x_race_B[i] + beta3*x_race_N[i] +
      beta4*x_race_A[i] + beta5*x_race_P[i] + beta6*x_race_M[i], invsigma2)
  }
  ## priors
  beta0 ~ dnorm(mu0, g0)
  beta1 ~ dnorm(mu1, g1)
  beta2 ~ dnorm(mu2, g2)
  beta3 ~ dnorm(mu3, g3)
```

```

beta4 ~ dnorm(mu4, g4)
beta5 ~ dnorm(mu5, g5)
beta6 ~ dnorm(mu6, g6)
invsigma2 ~ dgamma(a, b)
sigma <- sqrt(pow(invsigma2, -1))
}
"

```

```

y = as.vector(data$logInc)
x_rural = as.vector(data$Rural)
x_race_B = as.vector(data$Race_Black)
x_race_N = as.vector(data$Race_NA)
x_race_A = as.vector(data$Race_Asian)
x_race_P = as.vector(data$Race_PI)
x_race_M = as.vector(data$Race_M)
N = length(y)

```

```

the_data <- list("y" = y,
  "x_rural" = x_rural, "x_race_B" = x_race_B,
  "x_race_N" = x_race_N, "x_race_A" = x_race_A,
  "x_race_P" = x_race_P, "x_race_M" = x_race_M,
  "N" = N,
  "mu0" = 0, "g0" = 1, "mu1" = 0, "g1" = 1,
  "mu2" = 0, "g2" = 1, "mu3" = 0, "g3" = 1,
  "mu4" = 0, "g4" = 1, "mu5" = 0, "g5" = 1,
  "mu6" = 0, "g6" = 1,
  "a" = 1, "b" = 1)

```

```

initsfunction <- function(chain){
  .RNG.seed <- c(1,2)[chain]
  .RNG.name <- c("base::Super-Duper",
    "base::Wichmann-Hill")[chain]
  return(list(.RNG.seed=.RNG.seed,
    .RNG.name=.RNG.name))
}

```

```

posterior_MLR <- run.jags(modelString,
  n.chains = 1,
  data = the_data,
  monitor = c("beta0", "beta1", "beta2",
    "beta3", "beta4", "beta5",
    "beta6", "sigma"),
  adapt = 1000,
  burnin = 5000,
  sample = 5000,
  thin = 1,
  inits = initsfunction)

```

```
## Warning: Convergence cannot be assessed with only 1 chain
```

```
post <- as.mcmc(posterior_MLR)
```

```

synthesize <- function(X_rural, X_RB, X_RN, X_RA, X_RP, X_RM, index, n){
  mean_Y <- post[index, "beta0"] + X_rural * post[index, "beta1"] + X_RB * post[index, "beta2"] + X_RN * post[index, "beta3"] +
  synthetic_Y <- rnorm(n, mean_Y, post[index, "sigma"])
  data.frame(synthetic_Y, X_rural, X_RB, X_RN, X_RA, X_RP, X_RM)
}

```

```
n <- dim(data)[1]
```

```


Syndata <- synthesize(data$Rural, data$Race_Black, data$Race_NA, data$Race_Asian, data$Race_PI, data$Race_M)
names(Syndata) <- c("logInc", "Rural", "RB", "RN", "RA", "RP", "RM")

```

- (b) Make a scatter plot of the synthesized log(Income) against the original log(Income), and see what you find.

```
Synthesis <- cbind(data$logInc, Syndata$logInc)
Synthesis <- data.frame(Synthesis)
names(Synthesis) <- c("Orig", "Syn")

ggplot(Synthesis, aes(x=Orig, y=Syn)) + geom_point()
```



(b) Compare the mean and median of $\log(\text{Income})$, in the original dataset and the confidential dataset. Are they close to each other?

```
mean(Synthesis$Orig)

## [1] 10.59507
median(Synthesis$Orig)

## [1] 10.70574
mean(Synthesis$Syn)

## [1] 10.48264
median(Synthesis$Syn)

## [1] 10.47355
```

(c) Compare the point estimate of the regression coefficients of $\log(\text{Income})$ on $\log(\text{Expenditure})$, in the original dataset and the confidential dataset. Are they close to each other?

```
Synthesis <- cbind(Synthesis, log(data$Expenditure))
names(Synthesis) <- c("Orig", "Syn", "Ex")
model_orig <- lm(Orig ~ Ex, data=Synthesis)
```

```
model_syn <- lm(Syn ~ Ex, data=Synthesis)
summary(model_orig)
```

```
##
## Call:
## lm(formula = Orig ~ Ex, data = Synthesis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.9086 -0.4371  0.1069  0.6110  2.8823
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4.11123    0.30801   13.35  <2e-16 ***
## Ex           0.73814    0.03489   21.15  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.958 on 992 degrees of freedom
## Multiple R-squared:  0.3109, Adjusted R-squared:  0.3102
## F-statistic: 447.5 on 1 and 992 DF,  p-value: < 2.2e-16
summary(model_syn)
```

```
##
## Call:
## lm(formula = Syn ~ Ex, data = Synthesis)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.0097 -0.7894 -0.0123  0.8537  3.3229
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 10.79697    0.36640   29.468  <2e-16 ***
## Ex          -0.03578    0.04151   -0.862   0.389
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.14 on 992 degrees of freedom
## Multiple R-squared:  0.0007486, Adjusted R-squared: -0.0002587
## F-statistic: 0.7432 on 1 and 992 DF,  p-value: 0.3888
```

- (d) Evaluate the global utility measures of your synthesized log(Income) from your Bayesian synthesis model for the CE sample.

```
Oridata <- cbind(data$logInc, data$Rural, data$Race_Black, data$Race_NA, data$Race_Asian, data$Race_PI, data$Race_HI)
Oridata <- data.frame(Oridata)
names(Oridata) <- c("logInc", "Rural", "RB", "RN", "RA", "RP", "RM")
merged_data <- rbind(Oridata, Syndata)
merged_data$T <- c(rep(0, 994), rep(1, 994))
```

I. Propensity Score Measure

```
log_reg <- glm(T ~ logInc + as.factor(Rural) + as.factor(RB) + as.factor(RN) + as.factor(RA) + as.factor(RP) + as.factor(RM), data=merged_data)
summary(log_reg)
```

```
##
## Call:
## glm(formula = T ~ logInc + as.factor(Rural) + as.factor(RB) +
##      as.factor(RN) + as.factor(RA) + as.factor(RP) + as.factor(RM),
##      family = "binomial", data = merged_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.34866  -1.17247  -0.02424   1.17506   1.29890
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      0.93377    0.42512   2.196  0.0281 *
## logInc          -0.08809    0.03981  -2.213  0.0269 *
## as.factor(Rural)1 -0.01195    0.20501  -0.058  0.9535
## as.factor(RB)1    -0.03967    0.14572  -0.272  0.7854
## as.factor(RN)1    -0.09123    0.54092  -0.169  0.8661
## as.factor(RA)1     0.02868    0.23266   0.123  0.9019
## as.factor(RP)1    -0.02488    0.58003  -0.043  0.9658
## as.factor(RM)1    -0.04222    0.34759  -0.121  0.9033
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2756  on 1987  degrees of freedom
## Residual deviance: 2751  on 1980  degrees of freedom
## AIC: 2767
##
## Number of Fisher Scoring iterations: 3
pred <- predict(log_reg, data=merged_data)
probs <- exp(pred)/(1+exp(pred))
Up <- 1/(2*n)*sum((probs-1/2)^2)
Up
```

```
## [1] 0.0006173112
```

II. Cluster Analysis Measure

```
clusters <- hclust(dist(merged_data[,1:2]), method='average')

ClusterAnalysis <- function(G){
  clusterCut <- cutree(clusters, G)
  cluster_T <- as.data.frame(cbind(clusterCut, merged_data$T))
  names(cluster_T) <- c("cluster", "T")
  table(cluster_T)
  n_gS <- table(cluster_T)[,1]
  n_g <- rowSums(table(cluster_T))
  w_g <- n_g/(2*n)
  Uc <- (1/G)*sum(w_g*(n_gS/n_g-1/2)^2)
  Uc
}
```

```
ClusterAnalysis(3)
```

```
## [1] 0.0001895369
```

```
ClusterAnalysis(5)
```

```
## [1] 0.0006810602
```

```
ClusterAnalysis(10)
```

```
## [1] 0.0004248599
```

```
ClusterAnalysis(15)
```

```
## [1] 0.0003931866
```

```
ClusterAnalysis(20)
```

```
## [1] 0.0003456766
```

III. Empirical CDF Measure

```
ecdf_ori <- ecdf(Oridata[, "logInc"])
ecdf_syn <- ecdf(Syndata[, "logInc"])
percentile_ori <- ecdf_ori(merged_data[, "logInc"])
percentile_syn <- ecdf_syn(merged_data[, "logInc"])
```

```
ecdf_diff <- percentile_ori - percentile_syn
Um <- max(abs(ecdf_diff))
Um
```

```
## [1] 0.1096579
```

```
Ua <- mean(ecdf_diff^2)
Ua
```

```
## [1] 0.003470492
```

Drechsler (2001)

- i. Generate $m = 20$ synthetic datasets given your synthesis model for the CE sample. If you are using `set.seed()`, make sure that you do not generate the same synthetic data for each $m = 20$.

```
Syn_sets <- NULL
n <- dim(data)[1]
Syn_sets <- cbind(data$Rural, data$Race_Black, data$Race_NA, data$Race_Asian, data$Race_PI, data$Race_M)
Syn_sets <- data.frame(Syn_sets)
names(Syn_sets) <- c("Rural", "RB", "RN", "RA", "RP", "RM")
synthesize <- function(X_rural, X_RB, X_RN, X_RA, X_RP, X_RM, index, n){
  mean_Y <- post[index, "beta0"] + X_rural * post[index, "beta1"] + X_RB * post[index, "beta2"] + X_RN
  synthetic_Y <- rnorm(n, mean_Y, post[index, "sigma"])
  data.frame(synthetic_Y)
}
```

```
for (i in 1:20) {
  Syndata <- synthesize(data$Rural, data$Race_Black, data$Race_NA, data$Race_Asian, data$Race_PI, data$Race_M)
  Syn_sets <- cbind(Syn_sets, Syndata$synthetic_Y)
  names(Syn_sets[, 6+i]) <- c(paste("syn", as.character(i)))
}
```

- ii. Estimate a few analysis-specific utility measures, e.g. the mean and median of a continuous synthetic variable, the regression analysis coefficients, for each synthetic dataset.

```

Syn_sets <- cbind(Syn_sets, data$logInc, log(data$Expenditure))
names(Syn_sets) <- c("Rural", "RB", "RN", "RA", "RP", "RM", "m1", "m2", "m3", "m4", "m5", "m6", "m7", "m8",
                    "m9", "m10", "m11", "m12", "m13", "m14", "m15", "m16", "m17", "m18", "m19", "m20", "ori", "ex")
mean <- c()
median <- c()
coeff <- c()
pVal <- c()
variance <- c()
for (i in 1:21) {
  Syn_sets$target <- Syn_sets[,i+6]
  mean<- append(mean, mean(Syn_sets$target))
  median <- append(median, median(Syn_sets$target))
  lr <- lm(target ~ ex, data=Syn_sets)
  coe <- summary(lr)$coefficients[2, 1]
  coeff <- append(coeff, coe)
  pvalue <- summary(lr)$coefficients[2, 4]
  pVal <- append(pVal, pvalue)
  variance <- append(variance, var(Syn_sets$target))
}
Analysis <- cbind(mean, median, coeff, pVal, variance)
Analysis <- data.frame(Analysis)
names(Analysis) <- c("mean", "median", "coeff", "pVal", "variance")
Analysis

```

##	mean	median	coeff	pVal	variance
## 1	10.51355	10.53726	-0.003141683	9.420850e-01	1.407583
## 2	10.59510	10.57208	0.107992662	8.583629e-03	1.275252
## 3	10.59988	10.62107	0.066399481	1.221123e-01	1.390165
## 4	10.51870	10.48205	0.034256341	4.334615e-01	1.440107
## 5	10.61464	10.59914	0.092065792	2.593074e-02	1.289319
## 6	10.63823	10.62265	0.097133955	2.443097e-02	1.406019
## 7	10.55714	10.52877	-0.006796671	8.750428e-01	1.406035
## 8	10.56683	10.54801	-0.025857912	5.529836e-01	1.430004
## 9	10.66747	10.70796	0.061097291	1.450718e-01	1.324671
## 10	10.52346	10.52383	0.050915946	2.319512e-01	1.366617
## 11	10.55542	10.53428	0.104186609	1.358523e-02	1.345442
## 12	10.59726	10.58756	0.027299539	5.352915e-01	1.459815
## 13	10.51637	10.48682	0.079748258	6.395211e-02	1.397287
## 14	10.60760	10.61100	-0.001927762	9.631433e-01	1.309966
## 15	10.54066	10.52085	0.063251254	1.344894e-01	1.345905
## 16	10.58647	10.59106	0.091814505	3.120686e-02	1.370222
## 17	10.51784	10.53160	0.026529956	5.382356e-01	1.398744
## 18	10.60544	10.57778	0.089748411	3.137324e-02	1.311820
## 19	10.56639	10.56163	0.101849275	9.158309e-03	1.153558
## 20	10.54352	10.52183	0.059238110	1.680769e-01	1.391479
## 21	10.59507	10.70574	0.738140359	2.849973e-82	1.330535

- iii. Use the combining rules in Drechsler 2001 Chapter 6-1 (for fully synthetic data) and / or Drechsler 2001 Chapter 7-1 (for partially synthetic data) and create your final point estimate and confidence interval of the analysis-specific utility measures you calculated in Item ii above.

For mean:

```

qm <- sum(Analysis[1:20,1])/20
bm <- sum((Analysis[1:20,1]-qm)^2)/(20-1)

```

```
um <- sum(Analysis[1:20,5])/20
```

```
qm
```

```
## [1] 10.5716
```

```
bm
```

```
## [1] 0.001939386
```

```
um
```

```
## [1] 1.361001
```

```
qm+1.645*bm
```

```
## [1] 10.57479
```

```
qm-1.645*bm
```

```
## [1] 10.56841
```

The final point estimate is then 10.59, with a 90% confidence interval of [10.588,10.594].

Drechsler, J. and Reiter, J. P. (2009)

- i. Calculate the corresponding interval overlap measure for each of the analysis-specific utility measures you have done in Item 2.ii above.

```
var_syn <- sum(Analysis[1:20,5])/20
```

```
var_ori <- sum(Analysis[21,5])
```

```
mean_syn <- sum(Analysis[1:20,1])/20
```

```
mean_ori <- sum(Analysis[21,1])
```

```
mean_syn + 1.645*var_syn
```

```
## [1] 12.81044
```

```
mean_syn - 1.645*var_syn
```

```
## [1] 8.332753
```

```
mean_ori + 1.645*var_ori
```

```
## [1] 12.7838
```

```
mean_ori - 1.645*var_ori
```

```
## [1] 8.406342
```

```
(12.81-8.37)/(2*(12.78-8.41))+(12.81-8.37)/(2*(12.81-8.37))
```

```
## [1] 1.008009
```