

analysis__2-10

Kevin Ros

2/6/2020

Collaboration: I wasn't sure where to start, and Henrik pointed me in the direction of last year's multiple linear regression slides: https://github.com/monika76five/Undergrad-Bayesian-Statistics/blob/master/Lectures/BayesianLinearRegression/F19MATH347_Regression_R.pdf which I implemented, as it wasn't too clear how to do hierarchical modeling with categorical variables (but I think fastDummies can be used). Additionally, Henrik and I discussed the simulation portion, namely how one needs to pass a data frame of each variable, and not just income.

```
library(ggplot2)
library(coda)
library(runjags)
library(fastDummies)
data = data.frame(read.csv("../datasets/CEdata.csv",header=TRUE))
```

```
# EXAMPLE FROM CLASS
```

```
# modelString <-"
# model {
# ## sampling
# for (i in 1:N){
# y[i] ~ dnorm(beta0 + beta1*x[i], invsigma2)
# }
#
# ## priors
# beta0 ~ dnorm(mu0, g0)
# beta1 ~ dnorm(mu1, g1)
# invsigma2 ~ dgamma(a,b)
# sigma <- sqrt(pow(invsigma2,-1))
# }
# "
#
# y <- as.vector(log(data$TotalIncomeLastYear))
# x <- as.vector(log(data$TotalExpLastQ))
# N <-length(y)
# the_data <- list("y" = y, "x" = x, "N" = N,
#                 "mu0" = 0, "g0" = 0.0001,
#                 "mu1" = 0, "g1" = 0.0001,
#                 "a" = 1, "b" = 1)
#
# initsfunction <- function(chain){
#   .RNG.seed <- c(1,2)[chain]
#   .RNG.name <- c("base::SuperDuper",
#                 "base::Wichmann-Hill")[chain]
#   return(list(.RNG.seed=.RNG.seed,
#               .RNG.name=.RNG.name))
# }
# ...
#
# ```{r}
```

```
# posterior <- run.jags(modelString,
#                       n.chains = 1,
#                       data = the_data,
#                       monitor = c("beta0", "beta1", "sigma"),
#                       adapt = 1000,
#                       burnin = 5000,
#                       sample = 5000,
#                       thin = 50)
# plot(posterior, vars = "beta0")
```

The original slides has scale(), but I removed it to better interpret the results.

```
data$log_TotalExpSTD <- log(data$Expenditure)
data$log_TotalIncomeSTD <- log(data$Income)
```

Create the binary columns and rows for each categorical variable:

```
data$Rural = fastDummies::dummy_cols(data$UrbanRural)[,names(fastDummies::dummy_cols(data$UrbanRural)) == ".data_1"]
data$Race_Black = fastDummies::dummy_cols(data$Race)[,names(fastDummies::dummy_cols(data$Race)) == ".data_2"]
data$Race_NA = fastDummies::dummy_cols(data$Race)[,names(fastDummies::dummy_cols(data$Race)) == ".data_3"]
data$Race_Asian = fastDummies::dummy_cols(data$Race)[,names(fastDummies::dummy_cols(data$Race)) == ".data_4"]
data$Race_PI = fastDummies::dummy_cols(data$Race)[,names(fastDummies::dummy_cols(data$Race)) == ".data_5"]
data$Race_M = fastDummies::dummy_cols(data$Race)[,names(fastDummies::dummy_cols(data$Race)) == ".data_6"]
```

Same parameters from slides:

```
modelString <- "
model {
  ## sampling
  for (i in 1:N){
    y[i] ~ dnorm(beta0 + beta1*x_income[i] + beta2*x_rural[i] +
    beta3*x_race_B[i] + beta4*x_race_N[i] +
    beta5*x_race_A[i] + beta6*x_race_P[i] +
    beta7*x_race_M[i], invsigma2)
  }
  ## priors
  beta0 ~ dnorm(mu0, g0)
  beta1 ~ dnorm(mu1, g1)
  beta2 ~ dnorm(mu2, g2)
  beta3 ~ dnorm(mu3, g3)
  beta4 ~ dnorm(mu4, g4)
  beta5 ~ dnorm(mu5, g5)
  beta6 ~ dnorm(mu6, g6)
  beta7 ~ dnorm(mu7, g7)
  invsigma2 ~ dgamma(a, b)
  sigma <- sqrt(pow(invsigma2, -1))
}
"
```

```
y = as.vector(data$log_TotalExpSTD)
x_income = as.vector(data$log_TotalIncomeSTD)
x_rural = as.vector(data$Rural)
x_race_B = as.vector(data$Race_Black)
x_race_N = as.vector(data$Race_NA)
x_race_A = as.vector(data$Race_Asian)
x_race_P = as.vector(data$Race_PI)
```

```
x_race_M = as.vector(data$Race_M)
N = length(y)
```

Same parameter values from slides:

```
the_data <- list("y" = y, "x_income" = x_income,
  "x_rural" = x_rural, "x_race_B" = x_race_B,
  "x_race_N" = x_race_N, "x_race_A" = x_race_A,
  "x_race_P" = x_race_P, "x_race_M" = x_race_M,
  "N" = N,
  "mu0" = 0, "g0" = 1, "mu1" = 0, "g1" = 1,
  "mu2" = 0, "g2" = 1, "mu3" = 0, "g3" = 1,
  "mu4" = 0, "g4" = 1, "mu5" = 0, "g5" = 1,
  "mu6" = 0, "g6" = 1, "mu7" = 0, "g7" = 1,
  "a" = 1, "b" = 1)
```

```
initsfunction <- function(chain){
  .RNG.seed <- c(1,2)[chain]
  .RNG.name <- c("base::Super-Duper",
    "base::Wichmann-Hill")[chain]
  return(list(.RNG.seed=.RNG.seed,
    .RNG.name=.RNG.name))
}
```

Thinning of 50 is needed, otherwise beta0 and beta1 have very high lag

```
posterior_MLR <- run.jags(modelString,
  n.chains = 1,
  data = the_data,
  monitor = c("beta0", "beta1", "beta2",
    "beta3", "beta4", "beta5",
    "beta6", "beta7", "sigma"),
  adapt = 1000,
  burnin = 5000,
  sample = 5000,
  thin = 50,
  inits = initsfunction)
```

```
## Loading required namespace: rjags
## Compiling rjags model...
## Calling the simulation using the rjags method...
## Note: the model did not require adaptation
## Burning in the model for 5000 iterations...
## Running the model for 250000 iterations...
## Simulation complete
## Calculating summary statistics...

## Warning: Convergence cannot be assessed with only 1 chain

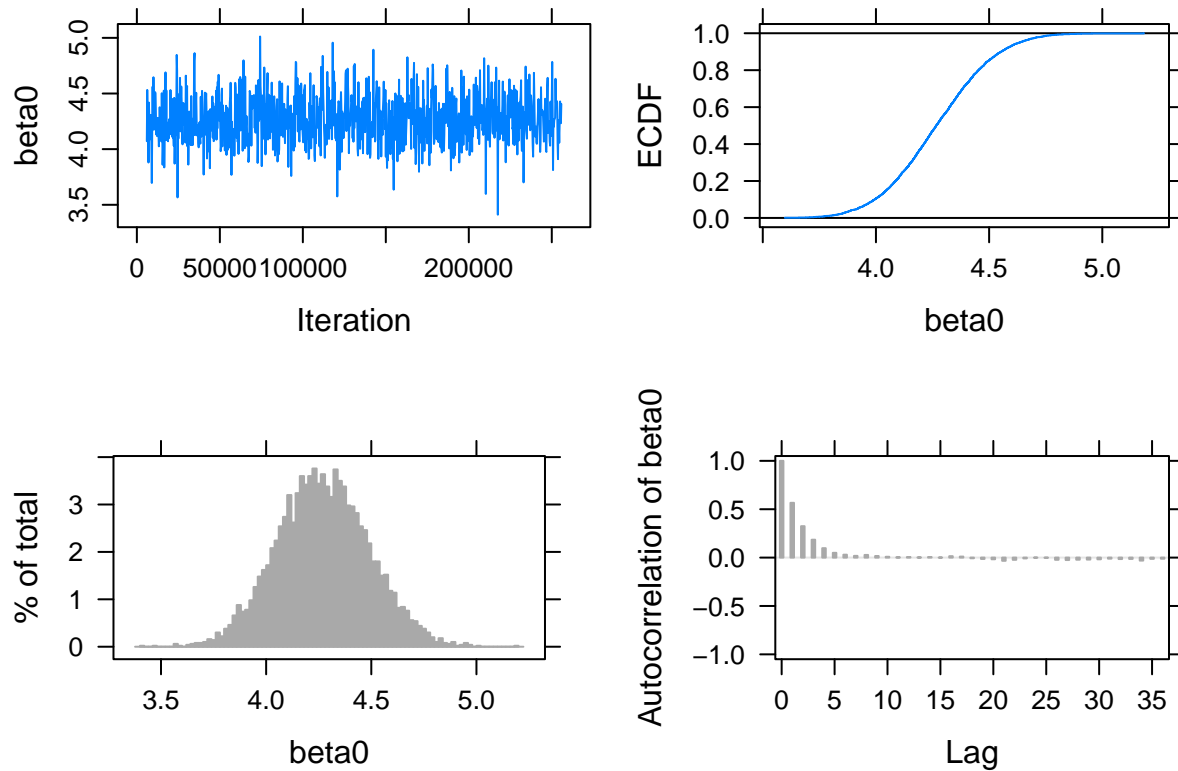
## Finished running the simulation
summary(posterior_MLR)
```









##	Lower95	Median	Upper95	Mean	SD	Mode
## beta0	3.8401288	4.265856654	4.68625457	4.270064662	0.21777079	NA
## beta1	0.3882419	0.428382748	0.46702186	0.428186927	0.02024773	NA
## beta2	-0.4324816	-0.227295748	-0.02303789	-0.226517253	0.10523080	NA









```
## beta3 -0.3306556 -0.190482478 -0.04140575 -0.191212993 0.07415444 NA
## beta4 -0.5206527 0.004840372 0.51719407 0.004623109 0.26616225 NA
## beta5 -0.0510838 0.162814921 0.40645797 0.162083182 0.11830143 NA
## beta6 -0.4595443 0.083013003 0.68580148 0.082703016 0.28996472 NA
## beta7 -0.2892273 0.040916578 0.39024470 0.042723529 0.17484430 NA
## sigma 0.6904107 0.721247329 0.75457051 0.721355983 0.01634697 NA
##
##          MCerr MC%ofSD SSeff          AC.500 psrf
## beta0 0.0058531144      2.7 1384 8.671215e-03 NA
## beta1 0.0005445194      2.7 1383 1.118156e-02 NA
## beta2 0.0014881883      1.4 5000 3.940920e-03 NA
## beta3 0.0010321668      1.4 5161 1.701377e-03 NA
## beta4 0.0037641027      1.4 5000 6.253260e-03 NA
## beta5 0.0017356791      1.5 4646 3.048365e-02 NA
## beta6 0.0041007204      1.4 5000 9.247818e-03 NA
## beta7 0.0024726717      1.4 5000 1.516543e-02 NA
## sigma 0.0002311811      1.4 5000 -1.238722e-05 NA
```









```
plot(posterior_MLR)
```









```
## Generating plots...
```















```
post <- as.mcmc(posterior_MLR)
```

```
synthesize <- function(X, index, n){
  mean_Y <- post[index, "beta0"] + X$x_income * post[index, "beta1"] + X$x_rural * post[index, "beta2"]
  synthetic_Y <- rnorm(n, mean_Y, post[index, "sigma"])
  data.frame(X$x_income, synthetic_Y)
}
```


```
n <- dim(data)[1]
params <- data.frame(x_income, x_rural, x_race_B, x_race_N, x_race_A, x_race_P, x_race_M)
synthetic_one <- synthesize(params, 1, n)
names(synthetic_one) <- c("OriginalIncome", "logIncome_syn")
```

Plotting the syntheticly generated income with the original income gives us the following:


```
plot(data$log_TotalIncomeSTD, synthetic_one$logIncome_syn)
```



```
plot(data$log_TotalExpSTD, data$log_TotalIncomeSTD)
```



```
plot(data$log_TotalExpSTD, synthetic_one$logIncome_syn)
```



The mean of the original income (10.59) is slightly more than the synthetic mean (8.77), a little more than one standard deviation.

```
mean(synthetic_one$OriginalIncome)
```

```
## [1] 10.59507
```

```
mean(synthetic_one$logIncome_syn)
```

```
## [1] 8.702602
```

Just as with the mean, the median of the original income (10.7) is slightly more than the median of the

synthetic income median (8.81), a little more than one standard deviation.

```
median(synthetic_one$OriginalIncome)
```

```
## [1] 10.70574
```

```
median(synthetic_one$logIncome_syn)
```

```
## [1] 8.751067
```

```
sd(synthetic_one$OriginalIncome)
```

```
## [1] 1.153488
```

```
sd(synthetic_one$logIncome_syn)
```

```
## [1] 0.9120538
```

Regarding the point estimates, with an expenditure of 0, original income will be 4.3219 and synthetic income will be 5.8609. And with every unit increase of expenditure, the original income increases by 0.4211 and the synthetic income increases by 0.3333. Of course, all of these values are in log form.

```
lm(data$log_TotalExpSTD ~ data$log_TotalIncomeSTD)
```

```
##
```

```
## Call:
```

```
## lm(formula = data$log_TotalExpSTD ~ data$log_TotalIncomeSTD)
```

```
##
```

```
## Coefficients:
```

```
##          (Intercept)  data$log_TotalIncomeSTD
```

```
##          4.3219          0.4211
```

```
lm(data$log_TotalExpSTD ~ synthetic_one$logIncome_syn)
```

```
##
```

```
## Call:
```

```
## lm(formula = data$log_TotalExpSTD ~ synthetic_one$logIncome_syn)
```

```
##
```

```
## Coefficients:
```

```
##          (Intercept)  synthetic_one$logIncome_syn
```

```
##          6.0300          0.3165
```