

SynthesisModel

Sarah Boese

2/6/2020

```
library(ProbBayes)
library(dplyr)
library(ggplot2)
require(gridExtra)
library(reshape)
library(runjags)
library(coda)
library(tidyverse)
library(fastDummies)
crcblue <- "#2905a1"
```

```
CESample <- read.csv("CESample2.csv")
```

I decided that I wanted to use all variables within CESampe logExpenditure, UrbanRural, Race) as estimators for logIncome. Thus, I used a Multilinear regression model in which I scaled Log Income and Log Expenditure by centering at 0 and dividing by standard deviation. I use the following MLR model (where * denotes a standardized continuous variable):

$$\begin{aligned} Y_i^* \mid \beta_0, \beta_1, \dots, \beta_7, \sigma, \mathbf{x}_i^* \stackrel{ind}{\sim} \text{Normal}(\beta_0 &+ \beta_1 x_{i,expenditure}^* + \beta_2 x_{i,rural}^* \\ &+ \beta_3 x_{i,race_B} + \beta_4 x_{i,race_N} \\ &+ \beta_5 x_{i,race_A} + \beta_6 x_{i,race_P} \\ &+ \beta_7 x_{i,race_M}, \sigma). \end{aligned} \quad (1)$$

```
CESample <- CESample %>%
  mutate(LogTotalIncome = log(TotalIncomeLastYear))
CESample <- CESample %>%
  mutate(LogTotalExp = log(TotalExpLastQ))
```

```
CESample$Log_TotalExpSTD <- scale(CESample$LogTotalExp)
CESample$Log_TotalIncomeSTD <- scale(CESample$LogTotalIncome)
## create indictor variable for Rural
CESample$Rural = fastDummies::dummy_cols(CESample$UrbanRural)[,names(fastDummies::dummy_cols(CESample$UrbanRural))
== ".data_2"]
```

```
## create indicator variables for Black (2), Native American (3),
## Asian (4), Pacific Islander (5), and Multi-race (6)
CESample$Race_Black = fastDummies::dummy_cols(CESample$Race)[,names(fastDummies::dummy_cols(CESample$Race))
CESample$Race_NA = fastDummies::dummy_cols(CESample$Race)[,names(fastDummies::dummy_cols(CESample$Race))
CESample$Race_Asian = fastDummies::dummy_cols(CESample$Race)[,names(fastDummies::dummy_cols(CESample$Race))
CESample$Race_PI = fastDummies::dummy_cols(CESample$Race)[,names(fastDummies::dummy_cols(CESample$Race))
CESample$Race_M = fastDummies::dummy_cols(CESample$Race)[,names(fastDummies::dummy_cols(CESample$Race))
```

```

modelString <- "
model {
  ## sampling
  for (i in 1:N){
    y[i] ~ dnorm(beta0 + beta1*x_exp[i] + beta2*x_rural[i] +
    beta3*x_race_B[i] + beta4*x_race_N[i] +
    beta5*x_race_A[i] + beta6*x_race_P[i] +
    beta7*x_race_M[i], invsigma2)
  }
  ## priors
  beta0 ~ dnorm(mu0, g0)
  beta1 ~ dnorm(mu1, g1)
  beta2 ~ dnorm(mu2, g2)
  beta3 ~ dnorm(mu3, g3)
  beta4 ~ dnorm(mu4, g4)
  beta5 ~ dnorm(mu5, g5)
  beta6 ~ dnorm(mu6, g6)
  beta7 ~ dnorm(mu7, g7)
  invsigma2 ~ dgamma(a, b)
  sigma <- sqrt(pow(invsigma2, -1))
}
"

```

- Pass the data and hyperparameter values to JAGS:

```

y = as.vector(CESample$Log_TotalIncomeSTD)
x_exp = as.vector(CESample$Log_TotalExpSTD)
x_rural = as.vector(CESample$Rural)
x_race_B = as.vector(CESample$Race_Black)
x_race_N = as.vector(CESample$Race_NA)
x_race_A = as.vector(CESample$Race_Asian)
x_race_P = as.vector(CESample$Race_PI)
x_race_M = as.vector(CESample$Race_M)
N = length(y) # Compute the number of observations

```

- Pass the data and hyperparameter values to JAGS:

```

the_data <- list("y" = y, "x_exp" = x_exp,
  "x_rural" = x_rural, "x_race_B" = x_race_B,
  "x_race_N" = x_race_N, "x_race_A" = x_race_A,
  "x_race_P" = x_race_P, "x_race_M" = x_race_M,
  "N" = N,
  "mu0" = 0, "g0" = 0.0001, "mu1" = 0, "g1" = 0.0001,
  "mu2" = 0, "g2" = 1, "mu3" = 0, "g3" = 1,
  "mu4" = 0, "g4" = 1, "mu5" = 0, "g5" = 1,
  "mu6" = 0, "g6" = 1, "mu7" = 0, "g7" = 1,
  "a" = 1, "b" = 1)

```

- Pass the data and hyperparameter values to JAGS:

```

initsfunction <- function(chain){
  .RNG.seed <- c(1,2)[chain]
  .RNG.name <- c("base::Super-Duper",
                 "base::Wichmann-Hill")[chain]
  return(list(.RNG.seed=.RNG.seed,
              .RNG.name=.RNG.name))
}

```

- Run the JAGS code for this model:

```

posterior_MLR <- run.jags(modelString,
  n.chains = 1,
  data = the_data,
  monitor = c("beta0", "beta1", "beta2",
              "beta3", "beta4", "beta5",
              "beta6", "beta7", "sigma"),
  adapt = 1000,
  burnin = 5000,
  sample = 5000,
  thin = 2,
  inits = initsfunction)

## Calling the simulation...
## Welcome to JAGS 4.3.0 on Tue Feb 11 12:42:06 2020
## JAGS is free software and comes with ABSOLUTELY NO WARRANTY
## Loading module: basemod: ok
## Loading module: bugs: ok
## . . Reading data file data.txt
## . Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 994
##   Unobserved stochastic nodes: 9
##   Total graph size: 9984
## . Reading parameter file inits1.txt
## . Initializing model
## . Adaptation skipped: model is not in adaptive mode.
## . Updating 5000
## -----| 5000
## ***** 100%
## . . . . . Updating 10000
## -----| 10000
## ***** 100%
## . . . . Updating 0
## . Deleting model
## .
## Note: the model did not require adaptation
## Simulation complete. Reading coda files...
## Coda files loaded successfully
## Calculating summary statistics...

## Warning: Convergence cannot be assessed with only 1 chain

```

```
## Finished running the simulation
```

JAGS output for the MLR model

```
summary(posterior_MLR)
```

##	Lower95	Median	Upper95	Mean	SD	Mode
## beta0	-0.0311177	0.02608835	0.0869957	0.02606177	0.02980286	NA
## beta1	0.4918340	0.54553450	0.5955600	0.54534824	0.02668857	NA
## beta2	-0.2743130	-0.04835345	0.1952340	-0.04751894	0.11885285	NA
## beta3	-0.3097090	-0.14332450	0.0261562	-0.14148264	0.08690774	NA
## beta4	-1.1226000	-0.51939250	0.0605914	-0.52345066	0.30013430	NA
## beta5	-0.1497080	0.11270250	0.3803880	0.11387820	0.13525668	NA
## beta6	-0.9134750	-0.28965850	0.3296200	-0.28791357	0.31720398	NA
## beta7	-0.8128310	-0.42725700	-0.0305582	-0.43054482	0.20023325	NA
## sigma	0.7932920	0.82832250	0.8669630	0.82850990	0.01870803	NA

##	MCerr	MC%ofSD	SEff	AC.20	psrf
## beta0	0.0004448655	1.5	4488	-0.002126058	NA
## beta1	0.0003774333	1.4	5000	-0.001473396	NA
## beta2	0.0016808331	1.4	5000	-0.009477521	NA
## beta3	0.0012297711	1.4	4994	-0.010830759	NA
## beta4	0.0039695622	1.3	5717	0.034892325	NA
## beta5	0.0019128183	1.4	5000	0.004351354	NA
## beta6	0.0044859417	1.4	5000	0.021879566	NA
## beta7	0.0027856122	1.4	5167	-0.015481990	NA
## sigma	0.0002733078	1.5	4685	0.001118300	NA

```
post_MLR <- as.mcmc(posterior_MLR)
```

```
synthesize <- function(X, index, n){
  synth_Y=vector(mode="numeric", length = n)
  for(i in 1:n){
    mean_Y <- post_MLR[index, "beta0"] + X[i,1] * post_MLR[index, "beta1"] + X[i,2] *post_MLR[index, "beta2"] +
      X[i,3]*post_MLR[index, "beta3"] + X[i,4] *post_MLR[index, "beta4"] + X[i,5] *post_MLR[index, "beta5"] +
      X[i,6]*post_MLR[index, "beta6"] + X[i,7] *post_MLR[index, "beta7"]
    synth_Y[i]<- rnorm(1, mean_Y, post_MLR[index, "sigma"])
  }
  synthetic_frame<-as.data.frame(X, row.names = NULL, optional = FALSE)
  synthetic_frame<-add_column(synthetic_frame, synth_Y)
  return(synthetic_frame)
}
```

```
n <- dim(CESample)[1]
matrix_of_X<-matrix(nrow=n, ncol=7)
matrix_of_X[,1]<-as.vector(CESample$Log_TotalExpSTD)
matrix_of_X[,2]<-as.vector(CESample$Rural)
matrix_of_X[,3]<-as.vector(CESample$Race_Black)
matrix_of_X[,4]<-as.vector(CESample$Race_NA)
matrix_of_X[,5]<-as.vector(CESample$Race_Asian)
matrix_of_X[,6]<-as.vector(CESample$Race_PI)
matrix_of_X[,7]<-as.vector(CESample$Race_M)
```


```
synthetic_new <- synthesize(matrix_of_X, 1, n)
names(synthetic_new) <- c("logExpenditure", "Rural", "Black", "Native American", "Asian", "Pacific Islander")
```

```
SyntheticData <- data.frame(synthetic_new$logIncome_syn, CESample$logTotalIncomeSTD)
names(SyntheticData) = c("SyntheticLogIncomeSTD", "LogTotalIncomeSTD")
```

```
data<- melt(SyntheticData)
```


```
## Using as id variables
```

```
ggplot(data,aes(x=value, fill=variable)) + geom_density(alpha=0.25)
```



Here we can see the effect our model has on the data. The distribution curve, while similar to the data, looks to follow a normal curve more faithfully.

```
ggplot(SyntheticData, aes(x=SyntheticLogIncomeSTD, y=LogTotalIncomeSTD)) +
  geom_point()
```



```
summary = summarize_all(SyntheticData, .funs=c(mean, median))
names(summary) = c("SyntheticLogIncomeMean", "SyntheticLogIncomeMedian", "LogTotalIncomeMean", "LogTotalIncomeMedian")
summary
```

```
## SyntheticLogIncomeMean SyntheticLogIncomeMedian LogTotalIncomeMean
## 1 -0.003130824 1.086035e-16 0.04843122
## LogTotalIncomeMedian
## 1 0.09593896
```

These look fairly close to each other to me.